

ADVENTURES IN MACHINE LEARNING

LEARN AND EXPLORE MACHINE LEARNING

ABOUT

CODING THE DEEP LEARNING REVOLUTION EBOOK

CONTACT

EBOOK / NEWSLETTER SIGN-UP

21
Shares

Python TensorFlow Tutorial – Build a Neural Network

⌚ April 8, 2017 🚩 Andy 🌐 Deep learning, Neural networks, TensorFlow ⏱ 19



Google's TensorFlow has been a hot topic in deep learning recently. The open source software, designed to allow efficient computation of data flow graphs, is especially

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

POPULAR TUTORIALS

Neural Networks Tutorial
Pathway to Deep Learning
Python TensorFlow Tutorial – Build a Neural Network



Convolutional Neural Networks Tutorial in TensorFlow



Keras tutorial – build a convolutional neural network in 11 lines



Word2Vec word embedding tutorial in Python and TensorFlow

CATEGORIES

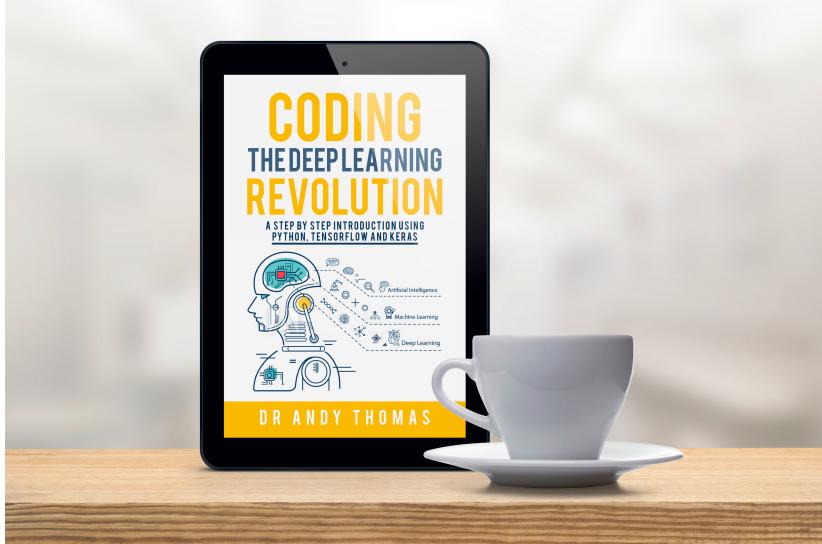
Amazon AWS

CNTK

Convolutional Neural Networks

incarnation – version 1.0 – it can even be run on certain mobile operating systems. This introductory tutorial to TensorFlow will give an overview of some of the basic concepts of TensorFlow in Python. These will be a good stepping stone to building more complex deep learning networks, such as **Convolution Neural Networks**, **natural language models** and **Recurrent Neural Networks** in the package. We'll be creating a simple three-layer neural network to classify the MNIST dataset. This tutorial assumes that you are familiar with the basics of neural networks, which you can get up to scratch with in the **neural networks tutorial** if required. To install TensorFlow, follow the instructions [here](#). The code for this tutorial can be found in [this site's GitHub repository](#). Once you're done, you also might want to check out a higher level deep learning library that sits on top of TensorFlow called Keras – see [my Keras tutorial](#).

Eager to learn more? Get the book here



First, let's have a look at the main ideas of TensorFlow.

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

[Ok](#)

gensim

GPUs

Keras

LSTMs

Neural networks

NLP

Optimisation

21
Shares

PyTorch

Recurrent neural network



Reinforcement learning

TensorBoard

TensorFlow



TensorFlow 2.0

Weight initialization

Word2Vec



NEWSLETTER + FREE EBOOK

Email address:

Your email address

[SIGN UP](#)

FIND US ON FACEBOOK

TensorFlow is based on graph based computation – “what on earth is that?”, you might say. It’s an alternative way of conceptualising mathematical calculations. Consider the following expression $a = (b + c) * (c + 2)$. We can break this function down into the following components:

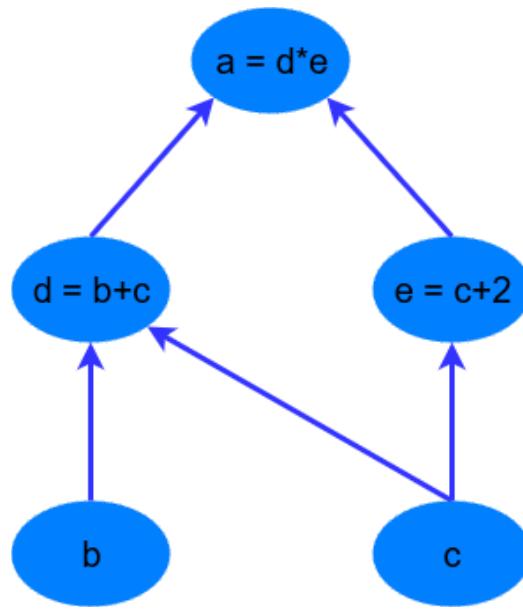
$$d = b + c$$

$$e = c + 2$$

$$a = d * e$$

Now we can represent these operations graphically as:

21
Shares



Simple computational graph

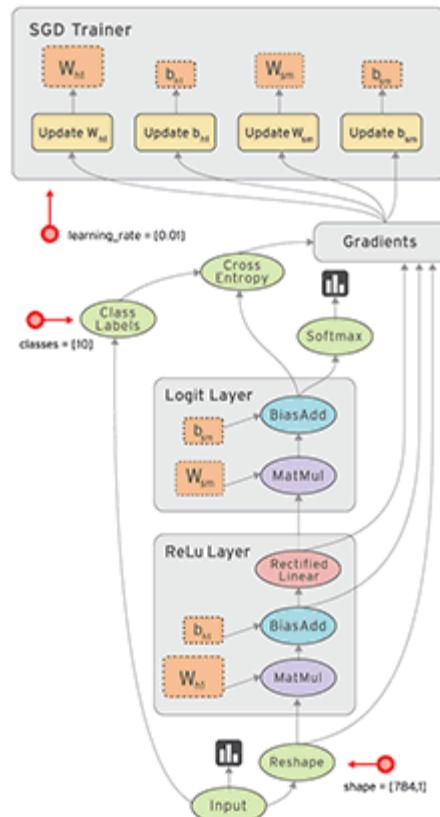
This may seem like a silly example – but notice a powerful idea in expressing the equation this way: two of the computations ($d = b + c$ and $e = c + 2$) can be performed in parallel. By splitting up these calculations across CPUs or GPUs, this can give us significant gains in computational times. These gains are a *must* for big data applications and deep learning – especially for complicated neural network architectures such as Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs). The idea

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

performance improvements via parallel operations and other efficiency gains.

We can look at a similar graph in TensorFlow below, which shows the computational graph of a three-layer neural network.



TensorFlow data flow graph

The animated data flows between different nodes in the graph are *tensors* which are multi-dimensional data arrays. For instance, the input data tensor may be $5000 \times 64 \times 1$, which represents a 64 node input layer with 5000 training samples. After the input layer there is a hidden layer with **rectified linear units** as the activation function. There is a final output layer (called a “**logit layer**” in the above graph).

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

block which finally flow to the **Stochastic Gradient Descent** optimiser which performs the back-propagation and gradient descent.

Here we can see how computational graphs can be used to represent the calculations in neural networks, and this, of course, is what TensorFlow excels at. Let's see how to perform some basic mathematical operations in TensorFlow to get a feel for how it all works.

21
Shares

2.0 A Simple TensorFlow example

Let's first make TensorFlow perform our little example calculation above - $a = (b + c) * (c + 2)$. First we need to introduce ourselves to TensorFlow *variables* and *constants*. Let's declare some then I'll explain the syntax:



```
1 import tensorflow as tf
2
3 # first, create a TensorFlow constant
4 const = tf.constant(2.0, name="const")
5
6 # create TensorFlow variables
7 b = tf.Variable(2.0, name='b')
8 c = tf.Variable(1.0, name='c')
```



As can be observed above, TensorFlow constants can be declared using the *tf.constant* function, and variables with the *tf.Variable* function. The first element in both is the value to be assigned the constant / variable when it is initialised. The second is an optional name string which can

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

/ variable from the initialised value, but it can also be set explicitly using the optional *dtype* argument. TensorFlow has many of its own types like *tf.float32*, *tf.int32* etc. – see them all [here](#).

It's important to note that, as the Python code runs through these commands, the variables haven't actually been declared as they would have been if you just had a standard Python declaration (i.e. *b* = 2.0). Instead, all the constants, variables, operations and the computational graph are only created when the initialisation commands are run.

21
Shares

Next, we create the TensorFlow *operations*:

```
1 # now create some operations
2 d = tf.add(b, c, name='d')
3 e = tf.add(c, const, name='e')
4 a = tf.multiply(d, e, name='a')
```



TensorFlow has a wealth of operations available to perform all sorts of interactions between variables, some of which we'll get to later in the tutorial. The operations above are pretty obvious, and they instantiate the operations $b + c$, $c + 2.0$ and $d * e$.



The next step is to setup an object to initialise the variables and the graph structure:

```
1 # setup the variable initialisation
2 init_op = tf.global_variables_initializer
```

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

session – *tf.Session*. The TensorFlow session is an object where all operations are run. TensorFlow was initially created in a *static* graph paradigm – in other words, first all the operations and variables are defined (the graph structure) and then these are compiled within the *tf.Session* object. There is now the option to build graphs on the fly using the TensorFlow Eager framework, to check this out see my [TensorFlow Eager tutorial](#).

However, there are still advantages in building static graphs using the *tf.Session* object. You can do this by using the *with* Python syntax, to run the graph like so:

21
Shares

```
1 # start the session
2 with tf.Session() as sess:
3     # initialise the variables
4     sess.run(init_op)
5     # compute the output of the graph
6     a_out = sess.run(a)
7     print("Variable a is {}".format(a_out))
```

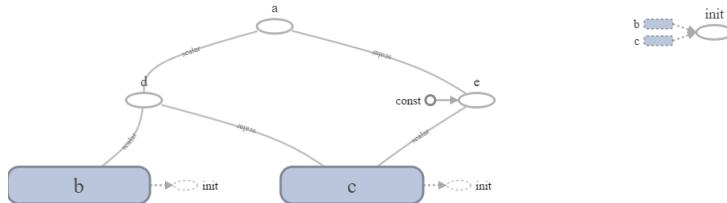
The first command within the *with* block is the initialisation, which is run with the, well, *run* command. Next we want to figure out what the variable *a* should be. All we have to do is run the operation which calculates *a* i.e. *a = tf.multiply(d, e, name='a')*. Note that *a* is an *operation*, not a variable and therefore it can be *run*. We do just that with the *sess.run(a)* command and assign the output to *a_out*, the value of which we then print out.

Note something cool – we defined operations *d* and *e* which need to be calculated before we can figure out what *a* is. However, we don't have to explicitly run *those* operations, as TensorFlow knows what other operations

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

through its data flow graph which shows it all the required dependencies. Using the TensorBoard functionality, we can see the graph that TensorFlow created in this little program:



21
Shares

Simple TensorFlow graph

Now that's obviously a trivial example – what if we had an array of b values that we wanted to calculate the value of a over?



2.1 The TensorFlow placeholder



Let's also say that we didn't know what the value of the array b would be during the declaration phase of the TensorFlow problem (i.e. before the `with tf.Session() as sess`) stage. In this case, TensorFlow requires us to declare the basic structure of the data by using the `tf.placeholder` variable declaration. Let's use it for b :



```
1 # create TensorFlow variables
2 b = tf.placeholder(tf.float32, [None, 1])
```

Because we aren't providing an initialisation in this declaration, we need to tell TensorFlow what data type each element within the *tensor* is going to be. In this case, we want to use `tf.float32`. The second argument is the shape of the data that will be "injected" into this variable. In this

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

variable (hence the "?"), the placeholder is willing to accept a *None* argument in the size declaration. Now we can inject as much 1-dimensional data that we want into the *b* variable.

The only other change we need to make to our program is in the *sess.run(a,...)* command:

```
1 | a_out = sess.run(a, feed_dict={b: np.array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])})
```

21
Shares

Note that we have added the *feed_dict* argument to the *sess.run(a,...)* command. Here we remove the mystery and specify exactly what the variable *b* is to be – a one-dimensional range from 0 to 10. As suggested by the argument name, *feed_dict*, the input to be supplied is a Python dictionary, with each key being the name of the *placeholder* that we are filling.

When we run the program again this time we get:

```
1 | Variable a is [[ 3.]
 2 |   [ 6.]
 3 |   [ 9.]
 4 |   [ 12.]
 5 |   [ 15.]
 6 |   [ 18.]
 7 |   [ 21.]
 8 |   [ 24.]
 9 |   [ 27.]
10 |   [ 30.]]
```



Notice how TensorFlow adapts naturally from a scalar

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

the data will flow through the graph.

Now we are ready to build a basic MNIST predicting neural network.

3.0 A Neural Network Example

Now we'll go through an example in TensorFlow of creating a simple three layer neural network. In future articles, we'll show how to build more complicated neural network structures such as convolution neural networks and recurrent neural networks. For this example though, we'll keep it simple. If you need to scrub up on your neural network basics, check out my [popular tutorial on the subject](#). In this example, we'll be using the MNIST dataset (and its associated loader) that the TensorFlow package provides. This MNIST dataset is a set of 28×28 pixel grayscale images which represent hand-written digits. It has 55,000 training rows, 10,000 testing rows and 5,000 validation rows.

21
Shares



We can load the data by running:

```
1 from tensorflow.examples.tutorials.mnist  
2 mnist = input_data.read_data_sets("MNIST
```

The `one_hot=True` argument specifies that instead of the labels associated with each image being the digit itself i.e. "4", it is a vector with "one hot" node and all the other nodes being zero i.e. [0, 0, 0, 1, 0, 0, 0, 0, 0]. This lets us easily feed it into the output layer of our neural network.

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

Next, we can set-up the placeholder variables for the training data (and some training parameters):

```

1 # Python optimisation variables
2 learning_rate = 0.5
3 epochs = 10
4 batch_size = 100
5
6 # declare the training data placeholders
7 # input x - for 28 x 28 pixels = 784
8 x = tf.placeholder(tf.float32, [None, 784])
9 # now declare the output data placeholder
10 y = tf.placeholder(tf.float32, [None, 10])

```

21
Shares

Notice the x input layer is 784 nodes corresponding to the $28 \times 28 (=784)$ pixels, and the y output layer is 10 nodes corresponding to the 10 possible digits. Again, the size of x is $(? \times 784)$, where the $?$ stands for an as yet unspecified number of samples to be input – this is the function of the *placeholder* variable.

Now we need to setup the weight and bias variables for the three layer neural network. There are always $L-1$ number of weights/bias tensors, where L is the number of layers. So in this case, we need to setup two tensors for each:

```

1 # now declare the weights connecting the
2 W1 = tf.Variable(tf.random_normal([784, 300]))
3 b1 = tf.Variable(tf.random_normal([300]))
4 # and the weights connecting the hidden
5 W2 = tf.Variable(tf.random_normal([300, 10])),
6 b2 = tf.Variable(tf.random_normal([10])),

```

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

connections between the input and hidden layer. This neural network will have 300 nodes in the hidden layer, so the size of the weight tensor $W1$ is [784, 300]. We initialise the values of the weights using a random normal distribution with a mean of zero and a standard deviation of 0.03. TensorFlow has a replicated version of the **numpy random normal function**, which allows you to create a matrix of a given size populated with random samples drawn from a given distribution. Likewise, we create $W2$ and $b2$ variables to connect the hidden layer to the output layer of the neural network.

21
Shares

Next, we have to setup node inputs and activation functions of the hidden layer nodes:

```
1 # calculate the output of the hidden layer
2 hidden_out = tf.add(tf.matmul(x, W1), b1)
3 hidden_out = tf.nn.relu(hidden_out)
```

12
13
14

In the first line, we execute the standard matrix multiplication of the weights ($W1$) by the input vector x and we add the bias $b1$. The matrix multiplication is executed using the *tf.matmul* operation. Next, we finalise the *hidden_out* operation by applying a **rectified linear unit** activation function to the matrix multiplication plus bias. Note that TensorFlow has a rectified linear unit activation already setup for us, *tf.nn.relu*.



This is to execute the following equations, as detailed in the **neural networks tutorial**:

$$\begin{aligned} z^{(l+1)} &= W^{(l)}x + b^{(l)} \\ h^{(l+1)} &= f(z^{(l+1)}) \end{aligned}$$

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

```

1 # now calculate the hidden layer output
2 # output layer
3 y_ = tf.nn.softmax(tf.add(tf.matmul(hidd

```

Again we perform the weight multiplication with the output from the hidden layer (*hidden_out*) and add the bias, *b2*. In this case, we are going to use a **softmax activation** for the output layer – we can use the included TensorFlow softmax function *tf.nn.softmax*.

21
Shares

We also have to include a cost or loss function for the optimisation / backpropagation to work on. Here we'll use the cross entropy cost function, represented by:

$$J = -\frac{1}{m} \sum_{i=1}^m \sum_{j=1}^n y_j^{(i)} \log(y_{j-}^{(i)}) + (1-y_j^{(i)}) \log(1-y_{j-}^{(i)})$$

Where $y_j^{(i)}$ is the ith training label for output node *j*, $y_{j-}^{(i)}$ is the ith predicted label for output node *j*, *m* is the number of training / batch samples and *n* is the number . There are two operations occurring in the above equation. The first is the summation of the logarithmic products and additions *across all the output nodes*. The second is taking a mean of this summation *across all the training samples*. We can implement this cross entropy cost function in TensorFlow with the following code:

```

1 y_clipped = tf.clip_by_value(y_, 1e-10, 1.0)
2 cross_entropy = -tf.reduce_mean(tf.reduce_sum(
3                                     + (1 - y) * tf.log(y_clipped), 1))

```

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

between 1e-10 to 0.999999. This is to make sure that we never get a case where we have a $\log(0)$ operation occurring during training – this would return NaN and break the training process. The second line is the cross entropy calculation.

To perform this calculation, first we use TensorFlow's `tf.reduce_sum` function – this function basically takes the sum of a given axis of the tensor you supply. In this case, the tensor that is supplied is the element-wise cross-entropy calculation for a single node and training sample i.e.: $y_j^{(i)} \log(y_{j_clipped}^{(i)}) + (1-y_j^{(i)}) \log(1-y_{j_clipped}^{(i)})$. Remember that y and $y_clipped$ in the above calculation are $(m \times 10)$ tensors – therefore we need to perform the first sum over the second axis. This is specified using the `axis=1` argument, where "1" actually refers to the second axis when we have a zero-based indices system like Python.

21
Shares



After this operation, we have an $(m \times 1)$ tensor. To take the mean of this tensor and complete our cross entropy cost calculation (i.e. execute this part $\frac{1}{m} \sum_{i=1}^m$), we use TensorFlow's `tf.reduce_mean` function. This function simply takes the mean of whatever tensor you provide it. So now we have a cost function that we can use in the training process.

Let's setup the optimiser in TensorFlow:

```
1 # add an optimiser
2 optimiser = tf.train.GradientDescentOpti
```

Here we are just using the gradient descent optimiser provided by TensorFlow. We initialize it with a learning rate,

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

descent see [here](#) and [here](#)) and the **backpropagation** for you. How easy is that? TensorFlow has a library of popular neural network training optimisers, see [here](#).

Finally, before we move on to the main show, were we actually run the operations, let's setup the variable initialisation operation and an operation to measure the accuracy of our predictions:

21
Shares

```
1 # finally setup the initialisation opera
2 init_op = tf.global_variables_initializer()
3
4 # define an accuracy assessment operation
5 correct_prediction = tf.equal(tf.argmax(
6   accuracy = tf.reduce_mean(tf.cast(correct_
```



The correct prediction operation *correct_prediction* makes use of the TensorFlow *tf.equal* function which returns *True* or *False* depending on whether two arguments supplied to it are equal. The *tf.argmax* function is the same as the [numpy argmax function](#), which returns the index of the maximum value in a vector / tensor. Therefore, the *correct_prediction* operation returns a tensor of size ($m \times 1$) of *True* and *False* values designating whether the neural network has correctly predicted the digit. We then want to calculate the mean accuracy from this tensor – first we have to cast the type of the *correct_prediction* operation from a Boolean to a TensorFlow float in order to perform the *reduce_mean* operation. Once we've done that, we now have an *accuracy* operation ready to assess the performance of our neural network.

3.2 Setting up the training

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

the full code below, then talk through it:

```
1 # start the session
2 with tf.Session() as sess:
3     # initialise the variables
4     sess.run(init_op)
5     total_batch = int(len(mnist.train.labels) / batch_size)
6     for epoch in range(epochs):
7         avg_cost = 0
8         for i in range(total_batch):
9             batch_x, batch_y = mnist.train.next_batch(batch_size)
10            _, c = sess.run([optimiser, cost], feed_dict={x: batch_x, y: batch_y})
11            avg_cost += c / total_batch
12        print("Epoch:", (epoch + 1), "cost =", "{:.3f}".format(avg_cost))
13    print(sess.run(accuracy, feed_dict={x: batch_x, y: batch_y}))
```

21
Shares

Stepping through the lines above, the first couple relate to setting up the *with* statement and running the initialisation operation. The third line relates to our mini-batch training scheme that we are going to run for this neural network. If you want to know about mini-batch gradient descent, check out this [post](#). In the third line, we are calculating the number of batches to run through in each training epoch. After that, we loop through each training epoch and initialise an *avg_cost* variable to keep track of the average cross entropy cost for each epoch. The next line is where we extract a randomised batch of samples, *batch_x* and *batch_y*, from the MNIST training dataset. The TensorFlow provided MNIST dataset has a handy utility function, *next_batch*, that makes it easy to extract batches of data for training.

The following line is where we run two operations. Notice

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

will be performed. As such, we get two outputs, which we have assigned to the variables `_` and `c`. We don't really care too much about the output from the *optimiser* operation but we want to know the output from the *cross_entropy* operation – which we have assigned to the variable `c`. Note, we run the *optimiser* (and *cross_entropy*) operation on the batch samples. In the following line, we use `c` to calculate the average cost for the epoch.

Finally, we print out our progress in the average cost, and after the training is complete, we run the *accuracy* operation to print out the accuracy of our trained network on the test set. Running this program produces the following output:

21
Shares

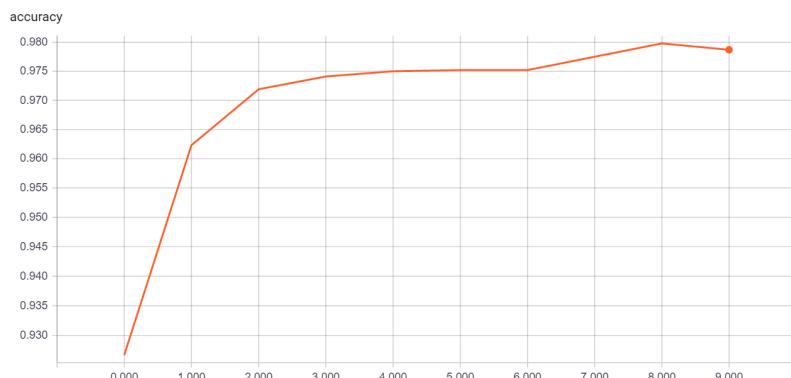
```
1 Epoch: 1 cost = 0.586
2 Epoch: 2 cost = 0.213
3 Epoch: 3 cost = 0.150
4 Epoch: 4 cost = 0.113
5 Epoch: 5 cost = 0.094
6 Epoch: 6 cost = 0.073
7 Epoch: 7 cost = 0.058
8 Epoch: 8 cost = 0.045
9 Epoch: 9 cost = 0.036
10 Epoch: 10 cost = 0.027
11
12 Training complete!
13 0.9787
```



There we go – approximately 98% accuracy on the test set, not bad. We could do a number of things to improve the model, such as regularisation (see this [tips and tricks post](#)), but here we are just interested in exploring TensorFlow. You can also use TensorBoard visualisation to look at things

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok

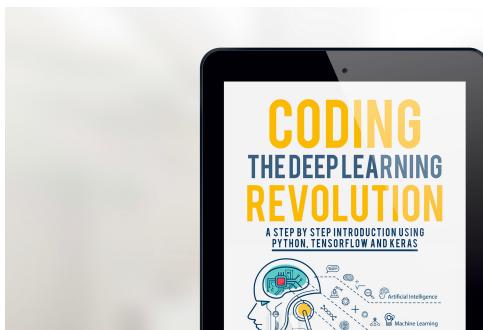
**TensorBoard plot of the increase in accuracy over 10 epochs**21
Shares

In a future article, I'll introduce you to TensorBoard visualisation, which is a really nice feature of TensorFlow. For now, I hope this tutorial was instructive and helps get you going on the TensorFlow journey. Just a reminder, you can check out the code for this post [here](#). I've also written an article that shows you how to build more complex neural networks such as [convolution neural networks](#), [recurrent neural networks](#) and [Word2Vec natural language models](#) in TensorFlow. You also might want to check out a higher level deep learning library that sits on top of TensorFlow called Keras – see [my Keras tutorial](#).

Have fun!

[ebook_store ebook_id="653"]

Eager to learn more? Get the book [here](#)



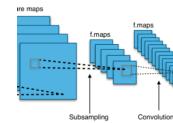
We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

**« PREVIOUS**

Stochastic
Gradient
Descent – Mini-
batch and more

NEXT »

Convolutional
Neural
Networks
Tutorial in
TensorFlow



21
Shares



Note: some posts
contain Udemy
affiliate links



Copyright © 2019 | WordPress Theme by MH Themes

We use cookies to ensure that we give you the best experience on our website. If you continue to use this site we will assume that you are happy with it.

Ok